

Extracting Biological Processes with Global Constraints

Author 1

XYZ Company
111 Anywhere Street
Mytown, NY 10000, USA
author1@xyz.org

Author 2

ABC University
900 Main Street
Ourcity, PQ, Canada A1A 1T2
author2@abc.ca

Abstract

Reasoning over processes is an important part of language understanding, and crucial for semantic applications such as questions answering. In this paper, we present the task of *process extraction*, in which process events and their relations are extracted from text. We first describe a novel graph representation for processes that includes a rich set of temporal, causal and co-reference relations between process events, and then characterize the structural properties of processes, such as *connectivity* of the process graph. We develop a method for predicting event-event relations in processes that exploits these structural properties by performing joint inference over the graph structure and rich relation set. We demonstrate empirically that taking process structure into account significantly improves performance over a data set containing 150 descriptions of biological processes. We release this novel data set for the benefit of the research community.

1 Introduction

Processes are complicated phenomena that involve a series of events and multiple participants. Automatically extracting the structure of processes from text is a necessary step for applications that require reasoning over process events. Consider, for example, the paragraph in Figure 1, which describes the biological process of ATP synthesis. A human reading this paragraph can answer questions such as:

1. *How do H^+ ions contribute to the production of ATP?*

2. *What causes the rotor to spin?*

3. *In ATP synthesis, what happens if the rotor fails to spin?*

These questions depend on extraction of the process structure and reasoning over causal and temporal relations between process events. Question answering systems that rely on bag-of-words representations will fail to correctly answer such questions.

Process extraction is related to two recent lines of work in Information Extraction – event extraction and timeline construction. The BioNLP 2009 and 2011 shared tasks (Kim et al., 2009; Kim et al., 2011) led to increasing interest in biomedical event extraction (Poon and Vanderwende, 2010; Miwa et al., 2010; Riedel and McCallum, 2011; McClosky et al., 2011; Björne et al., 2011), where given a single sentence annotated with protein mentions, events are identified and relations between events and proteins are extracted. In this shared task participants were asked to consider nine event types that are relevant for proteins (such as *Phosphorylation* and *Transcription*). Processes, on the other hand, are centered around discovering *relations* between events that span *multiple* sentences. In Figure 1 for instance, process extraction involves determining the relations between 8 events (*flowing*, *enter*, etc.) that are necessary for ATP production, which appear across four sentences. Note that the set of possible event types can not be restricted to a small closed set.

Timeline construction involves identifying temporal relations between events (Chambers and Jurafsky, 2008; Yoshikawa et al., 2009; Denis and Muller, 2011; Do et al., 2012; McClosky and Manning, 2012), and is thus related to process extrac-

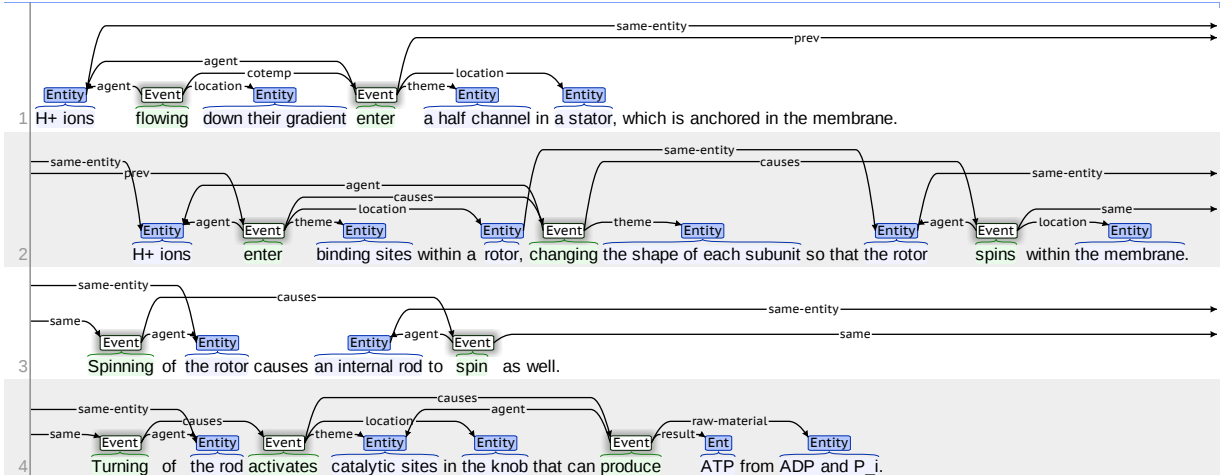


Figure 1: An annotation of the ATP synthesis process

tion as both focus on event-event relations that span multiple sentences. However, fully capturing process structure requires handling a rich set of relations such as CAUSES and SUPEREVENT (see Section 3), which are often not addressed in timeline construction. Moreover, processes exhibit particular properties that do not hold generally in temporal ordering. For example, in processes all events are somehow related to one another, a property that can be exploited for improving extraction.

In this paper, we present the task of process extraction and describe methods for extracting relations between process events. Our method works over multiple sentences and extracts a rich set of event-event relations, where the set of possible event types is open ended. Process structure is characterized by global properties that can be utilized during process extraction. For example, most processes exhibit a “chain-like” structure corresponding to process progression over time, and all process events are connected to one another, as previously noted. We will show that by incorporating global properties into our model and performing joint inference over the extracted relations we can significantly improve process quality. Our empirical experiments are performed over a novel data set of 150 process descriptions from the textbook “Biology” (Campbell and Reece, 2005) that were annotated by trained biologists. We note however that our method does not utilize any domain-specific knowledge and thus can be easily applied to domains other than Biology.

To conclude, this paper presents the following three contributions:

1. We define process extraction and characterize the structural properties of processes.
2. We show that modeling structural properties significantly improves process extraction comparing to several baselines.
3. We publicly release a novel data set of 150 fully annotated biological process descriptions.

2 Related Work

BioNLP work

Timeline construction work.

Scripts work - Chambers, Poon 2013.

Work that uses global constraints with ILP or dual decomposition or whatever.

3 Process Definition and Data Set

A process description is a paragraph or sequence of tokens $\mathbf{x} = \{x_1, \dots, x_{|\mathbf{x}|}\}$ describing a series of events that are related by various temporal and causal relations. For example, in ATP synthesis the event in which the rotor spins *causes* the event where an internal rod spins.

We define the process events and their relations by a directed graph $\mathcal{P} = (V, E)$, where the nodes $V = \{1, \dots, |V|\}$ represent event mentions and labeled edges correspond to event-event relations. An event mention $v \in V$ is defined by a trigger t_v , which is a span of words x_i, x_{i+1}, \dots, x_j and by

a set of argument mentions A_v , where each argument mention $a_v \in A_v$ is also a span of words labeled by a semantic role l taken from a set \mathcal{L} . For example, in the first event mention of ATP synthesis $t_v = \textit{flowing}$, and one of the arguments is $a_v = (\textit{H+ ions}, \text{AGENT})$. A labeled edge (u, v, r) in the graph describes a relation $r \in \mathcal{R}$ between the event mentions u and v . The task of process extraction is to extract the structure P from the text \mathbf{x} ¹.

A natural way to break down process extraction into two steps is to first perform semantic role labeling (SRL), that is, identify triggers and predict argument mentions with their semantic role, and then extract event-event relations between pairs of event mentions. In this paper, we focus on the second task, where given a set of triggers \mathcal{T} , we find all event-event relations. For completeness, we now describe the set of semantic roles \mathcal{L} used in our data set, and then present the set of event-event relations \mathcal{R} .

The set \mathcal{L} contains standard semantic roles such as AGENT, THEME, ORIGIN, DESTINATION and LOCATION. Two additional semantic roles were employed that are relevant for biological text: RESULT corresponds to an entity that is the result of an event, and RAW-MATERIAL describes an entity that is used or consumed during an event. For example, in the last event in Figure 1 ATP is the RESULT of the event, while ADP is the RAW-MATERIAL.

The relation set \mathcal{R} contains the following relations (assuming an edge (u, v, r)):

1. PREV denotes that u is an event immediately previous v . Thus, the edges (u, v, PREV) and (v, w, PREV) , preclude the edge (u, w, PREV) . For example, in “When a photon *strikes* ... energy is *passed* ... until it *reaches* ...”, there is no edge $(\textit{strikes}, \textit{reaches}, \text{PREV})$ due to the intervening event ‘*passed*’.
2. COTEMP denotes that events u and v overlap (e.g., the first two event mentions in Figure 1).
3. SUPER denotes that event v is included in event u . For instance, the process for “During *DNA replication*, DNA polymerases *proofread* each nucleotide...” has the edge $(\textit{DNA replication}, \textit{proofread}, \text{SUPER})$.

¹Argument mentions can also be related by coreference, but we neglect that since it is not central to this paper.

	Avg	Min	Max
# of tokens			
# of events			
# of relations			

Table 1: Statistics over the 150 process descriptions

4. CAUSES denotes that event u causes event v (e.g., the relation between *changing* and *spins* in sentence 2 of Figure 1).
5. ENABLES denotes that event u creates preconditions that allow event v to take place. For example, the description “... cause cancer cells to *lose* attachments to neighboring cells..., allowing them to *spread* into nearby tissues” has the edge $(\textit{lose}, \textit{spread}, \text{ENABLES})$.
6. SAME denotes that u and v co-refer to the same event (see Figure 1).

Our relation set contains the relations CAUSES and ENABLES, which are important for modeling processes and go beyond temporal ordering only. We defined that whenever these two relations apply they override the temporal relation (which is invariably PREV). The SUPER relation appears in temporal annotations such as The Timebank corpus (Pustejovsky et al., 2003) and in work on temporal logic (Allen, 1983), but in practice it is not considered by many temporal ordering systems (Chambers and Jurafsky, 2008; Yoshikawa et al., 2009; Do et al., 2012).

We also added event coreference (SAME) to \mathcal{R} . Do et al. (2012) used event coreference information in a temporal ordering task to modify probabilities provided by pairwise classifiers prior to joint inference. In this paper, we simply treat SAME as another event-event relation, which allows us to easily perform joint inference and employ structural constraints that combine both coreference and temporal relations simultaneously. For example, if (u, v, SAME) , then it can not be for any w that u is before w , but v is after w (see Section 4.3)

We have annotated 150 process descriptions based on the aforementioned definitions and provide further details on annotation and data set statistics in Section 5.4 and Table 1.

Structural properties of processes Naturally, coherent processes exhibit many structural properties.

Deg.	Gold	Local	Global
0			
1			
2			
3			

Table 2: Node degree count for event mentions across the process descriptions

For example, two argument mentions related to the same event mention can not overlap – a constraint that has been used in the past in SRL (Toutanova et al., 2008). In this paper we focus on three main structural properties of the graph \mathcal{P} . First, in a coherent process all event mentioned are related to one another, and hence the graph \mathcal{P} must be connected. Second, processes tend to have a “chain-like” structure where one event follows another. Thus, we expect node degree to generally be ≤ 2 , and this is indeed the case as demonstrated by the first column in Table 2. Last, if we consider all possible relation triangles, clearly some triangles are impossible, while other are common, which is illustrated in Figure ??.

In Section 4.3, we will show how using these properties we can improve process extraction, by formulating the problem as an ILP with both hard and soft constraints, and performing joint inference.

4 Joint Model for Process Extraction

Given a paragraph \mathbf{x} and a trigger set \mathcal{T} we wish to extract all event-event relations E . Similar to Do et al. (2012) our model consists of a local pairwise classifier and global constraints. We first introduce a classifier that is based on features from previous work (Section 4.1). Next, we describe novel features specific for process extraction (Section 4.2). Last, we incorporate global constraints into our model in an ILP formulation (Section 4.3).

4.1 Local pairwise classifier

The pairwise classifier predicts relations between all event mention pairs (represented by their triggers). Since relations in \mathcal{R} are directed, we must predict also the direction of each relation. We do this by expanding \mathcal{R} to include reverse relations and so we re-define \mathcal{R} to include 11 relations: PREV, NEXT, COTEMP, SUPER, SUB, CAUSES, CAUSED, ENABLES, ENABLED, SAME, NONE, where NONE in-

Feature	Description
POS	Pair of POS tags
Lemma	Pair of lemmas
Prep*	Preposition lexeme, if in a prepositional phrase
Words between	For adjacent triggers, content words between triggers
Temp. between	For adjacent triggers, temporal connectives (from a small list) between triggers
# sent.	Quantized number of sentences between triggers
# word.	Quantized number of words between triggers
LCA	Least common ancestor on constituency tree, if exists
Dominates*	Whether one trigger dominates other
Share	Whether triggers share a child on dependency tree

Table 3: Features extracted for a trigger pair (t_i, t_j) . Asteriks (*) indicate features that are duplicated, once for each trigger.

icates no relation. Thus, our classifier is a function $f : \mathcal{T} \times \mathcal{T} \rightarrow \mathcal{R}$. Let n be the number of triggers in a process description, and t_i be the i 'th trigger appearing in the description, since $f(t_i, t_j)$ completely determines $f(t_j, t_i)$ it suffices to consider only pairs such that $i < j$. Note that in this new definition of \mathcal{R} the process graph $\mathcal{P} = (V, E)$ is undirected.

Table 3 describes features from previous work (Chambers and Jurafsky, 2008; Do et al., 2012) extracted for a trigger pair (t_i, t_j) . Some features were omitted since they did not yield improvement in performance on a development set, or they require gold annotations provided in TimeBank, which we do not have. To reduce sparseness, we convert nominalizations into their verbal forms when computing word lemmas, using WordNet's (Fellbaum, 1998) derivation links.

4.2 Classifier extensions

A central source of information for extracting event-event relations from text are *connectives* such as *after*, *during*, etc. However, there is variability in the

occurrence of these connectives. Consider the following two sentences (connectives in bold, triggers in *italics*):

1. **Because** alleles are *exchanged* during *gene flow*, genetic differences are *reduced*.
2. During *gene flow*, alleles are *exchanged*, and **so** genetic differences are *reduced*.

Both sentences express the relation (*exchanged*, *reduced*, CAUSES), but the connective used is different, its linear position with respect to the triggers is different, and in sentence 1 the trigger *gene flow* intervenes between *exchanged* and *reduced*. Since our data set is very small, we would like to identify the triggers related to each connective, and share features between such sentences. We do this using the dependency structure and a clustering of connectives.

[Details on clustering and syntactic feature implementation probably using example above, maybe explain using "marker", then say something that we do something similar for things that are under PP but do not describe for simplicity, and say that for advmod we only use the clustering]

We further extend our features to handle the rich relation set necessary for process extraction. Processes often begin with a trigger that is a SUPER of the subsequent triggers, e.g., "The *Calvin cycle* begins by *incorporating*...". Thus, we add a feature for t_i indicating whether $i = 1$ and t_i is a noun. We also add two features targeted at the relation SAME: one indicating whether the lemmas of t_i and t_j are equal, and another specifying the determiner of t_j , if it exists. Intuition is that certain determiners indicate that the event triggered had already been mentioned, e.g., the determiner *this* hints a SAME relation in "The next steps *decompose* citrate back to oxaloacetate. This *regeneration* makes...". Last, we add as a feature the dependency path between t_i and t_j , if it exists, e.g., the feature $\xrightarrow{dobj} \xrightarrow{rmod}$ between *produces* and *divide* will fire in "meiosis produces cells that divide...".

For our pairwise classifier, we train a MaxEnt classifier that provides a probability p_{ijr} for every trigger pair (t_i, t_j) and relation r . Hence, $f(t_i, t_j) = \arg \max_r p_{ijr}$.

4.3 Global Constraints

Naturally, a pairwise classifier can result in a process structure that violates global properties as described in Section 3. Figure ?? shows in black edges the predictions of our local classifier, which result in the trigger [it'd be good to find a structure that combines transitivity violation and triad violation and refer to it] In this section we incorporate constraints that result in a coherent global process structure.

Let θ_{ijr} score the relation r for the triggers (t_i, t_j) (e.g, $\theta_{ijr} = \log p_{ijr}$), and y_{ijr} be an indicator for r . Our goal is to find an assignment for the indicators $\mathbf{y} = \{y_{ijr} \mid 1 \leq i < j \leq n, r \in \mathcal{R}\}$. With no global constraints this can be formulated with the following ILP:

$$\begin{aligned} \arg \max_{\mathbf{y}} \quad & \sum_{ijr} \theta_{ijr} y_{ijr} \\ \text{s.t.} \quad & \forall i, j \sum_r y_{ijr} = 1 \end{aligned} \quad (1)$$

where the constraint ensures each trigger pair is assigned to exactly one relation. We now describe constraints that result in a process that is connected, "chain-like", and contains reasonable relation triangles.

Connectivity Our formulation for enforcing connectivity is a minor variation on the one suggested by Martins et al. (2009) for dependency parsing. In our setup, we want \mathcal{P} to be an undirected connected graph, and not a directed tree. However, an undirected graph is connected iff there is a directed tree embedded in it, and thus the resulting formulation is almost identical. This formulation is based on flow constraints that ensure that there is a path from a root in the graph to all other nodes.

Let $\bar{\mathcal{R}}$ be the set of relations except for NONE, i.e., $\mathcal{R} \setminus \text{NONE}$. An edge (t_i, t_j) is in E if $y_{ij} = \sum_{r \in \bar{\mathcal{R}}} y_{ijr} = 1$. For each variable y_{ij} we define two auxiliary binary variables z_{ij} and z_{ji} that correspond to the directed edges in the embedded tree. We tie each auxiliary variable to its corresponding ILP variable by ensuring that it is active only if the ILP variable is active:

$$\forall i < j \quad z_{ij} < y_{ij}, z_{ji} < y_{ji} \quad (2)$$

Next, we add constraints that enforce the structure of the auxiliary variables results in a tree rooted in node 1. For that we need for every $i \neq j$ a flow variable ϕ_{ij} which specifies the amount of flow on the directed edge z_{ij} .

$$\sum_i z_{i1} = 0, \forall_{j \neq 1} \sum_i z_{ij} = 1 \quad (3)$$

$$\sum_i \phi_{1i} = n - 1 \quad (4)$$

$$\forall_j \sum_i \phi_{ij} - \sum_k \phi_{jk} = 1 \quad (5)$$

$$\forall_{i \neq j} \phi_{ij} \leq n \cdot z_{ij} \quad (6)$$

Equation 3 says that all nodes in the graph have exactly one parent, except for the root that has no parents. Equation 4 ensures that the outgoing flow from the root is $n-1$, and Equation 5 states that each of the other $n-1$ nodes consumes exactly one flow unit. Last, Equation 6 ties the auxiliary variables to the flow variables, making sure that flow occurs only on edges. The combination of these constraints guarantees that the graph induced by the variables z_{ij} is a directed tree and consequently the graph induced by our variables y is connected.

Chain structure A connected graph where the degree of all nodes is ≤ 2 is a chain. Table 2 presents nodes' degree and demonstrated that indeed process graphs are close to being chains. The following constraint bounds nodes' degree by 2:

$$\forall_j \sum_{i < j} y_{ij} + \sum_{j < k} y_{jk} \leq 2 \quad (7)$$

Since graph structures are not always chains we add this as a soft constraint, that is, we penalize the objective for each violation. Thus, our modified objective function is $\sum_{i,j,r} \theta_{ijr} y_{ijr} + \sum_{k \in \mathcal{K}} \alpha_k C_k$, where \mathcal{K} is the set of soft constraints, α_k is the penalty, and C_k indicates whether a constraint is violated. We tune the parameters α_k on a development set, as explained in Section 5.4.

Relation triangles A triangle is a 3-tuple of relations $(f(t_i, t_j), f(t_j, t_k), f(t_i, t_k))$. Clearly, some triangles are impossible while others are quite common. In order to look for triangles that could potentially improve process extraction we counted how

many time each possible triangle occurs in both the training data and the output of our pairwise classifier, and focused on those for which the classifier and the gold standard disagree. We are interested in triangles that never occur in the training data but are predicted by the classifier, and triangles that are frequent in the gold standard but do not appear in the classifier output. Figure ?? illustrates the triangles found and Equations X-Y provide the corresponding ILP formulation. Soft constraints are incorporated by defining a penalty/reward α_k for each triangle structure and expanding the set \mathcal{K} accordingly.

1. **SAME transitivity:** The transitivity of SAME (co-reference) has been used in past work (Finkel and Manning, 2008) and we incorporate it as a soft constraint (Equation X).
2. **COTEMP transitivity:** If t_i is co-temporal with t_j and t_j is co-temporal with t_k , then usually t_i and t_k are either co-temporal or denote the same event (soft constraint in Equation X).
3. **CAUSE-COTEMP:** If t_i causes both t_j and t_k , then often t_j and t_k are co-temporal. E.g, in “*genetic drift* has led to a *loss* of genetic variation and an *increase* in the frequency of harmful alleles”, a single event causes two subsequent events that occur simultaneously (soft constraint in Equation X).
4. **SAME contradiction:** if t_i is the same event as t_j , then their temporal ordering with respect to a third trigger t_k may result in a contradiction, e.g., if t_i is before t_k , but t_j is after t_k . We define 5 temporal categories that generate $\binom{5}{2}$ possible contradictions, but for brevity present just one representative hard constraint (Equation X). Note that this constraint depends on co-reference and temporal relations being predicted jointly.
5. **PREV:** As mentioned (Section 4.3), if t_i is immediately before t_j , and t_j is immediately before t_k , then it can not be that t_i is immediately before t_k (hard constraint in Equation X).

We used the Gurobi optimization package³ to find an exact solution for our ILP, which contains $O(n^2 r)$

³www.gurobi.com

variables and $O(n^3)$ constraints. We have also developed an equivalent formulation amenable to dual decomposition (Reichart and Barzilay, 2012), which is a faster approximation method, but practically we found that solving the problem exactly with Gurobi is quite fast (average/median time per process: XX/YY).

5 Experimental Evaluation

5.1 Experimental setup

Our data set consists of 148 process descriptions annotated by a biologist. The biologist was presented with annotation guidelines, annotated 20 descriptions and then the annotations were discussed with the authors, after which all process descriptions were annotated. We trained a second biologist on 20 process descriptions, and measured inter-annotator agreement on another 30 random process descriptions, resulting in good agreement $\kappa = 0.75$.

Process descriptions were parsed with the Stanford constituency and dependency parsers (Klein and Manning, 2003; de Marneffe et al., 2006), and 35 process descriptions were set aside as a test set (# of training set trigger pairs: XXX, # of test set trigger pairs: XXX). We performed 10-fold cross validation over the training set for feature selection and tuning of constraint parameters. For each constraint type (connectivity, chain-structure, and five triangle constraints) we introduced a parameter and tuned the seven parameters by coordinate-wise ascent, where for hard constraints a binary parameter controls whether the constraint is used, and for soft constraints it controls reward/penalty (α_k).

We test the following systems: (a) *All-Prev*: since the most common process structure is a chain of consecutive events we simply predict NEXTEVENT for every two adjacent triggers. (b) *Local_{base}*: A pairwise classifier with features from previous work (Section 4.1) (c) *Local*: A pairwise classifier with all features (Section 4.2) (d) *Local_{chain}*: For every two adjacent triggers we pick the highest probability non-NONE relation using the local classifier. Again, this uses the assumption that many processes have a chain-structure. (e) *Global*: Our full model that uses ILP inference.

To evaluate a system we compare its set of predictions on all trigger pairs to the gold standard

	Full			Temporal		
	P	R	F ₁	P	R	F ₁
<i>All-Prev</i>	34.1	32.0	33.0	62.2	58.3	60.2
<i>Local_{base}</i>						
<i>Local</i>	57.6	44.7	50.3	67.8	52.6	59.3
<i>Local_{chain}</i>	52.8	49.6	51.1	65.0	60.1	62.9
<i>Global</i>	55.4	54.0	54.7	66.2	64.5	65.3

Table 4: Test set results on all experiments

annotations and compute micro-averaged precision, recall and F₁. We perform two types of evaluations: (a) *Full*: evaluation on our full set of 11 relations (b) *Temporal*: Evaluation on temporal relations only, performed by collapsing PREVIOUS-EVENT, CAUSES, and ENABLES to a single category and similarly for NEXTEVENT, CAUSED, and ENABLED.

5.2 Results

Table 4

Discuss results:

- Global improves over local by 4.5 F1 points in the full evaluation (6 points in temporal)
- Local chain uses a fixed structure and uses the classifier to predict which relation - global is better in both precision and recall
- All-prev full evaluation is really bad but temporal evaluation is not bad - this is since indeed the structure is mostly chain-like with the two most common relations being PrevEvent and Causes. Still global is better in both precision and recall.
- our new features are important for our task they substantially improve the local classifier
- local classifier has good recall but bad precision - global improves recall thanks to connectivity constraint. Here we can talk about the double penalization.
- Mention that F₁ of human to human - that was not on the test set but shows that there is an upper bound around 0.7-0.75 for full evaluation and 0.85 for temporal

Table ?? is a confusion matrix showing that it is hard to distinguish Causes, PrevEvent, and Cotem-poral. These are not handled by global constraints and to handle this need other types of information.

5.3 Analysis and Discussion

Table that shows the setting of the hyper parameters on the dev set - what order and what are the parameter values - this shows what constraint help and what constraints don't. We can interpret that - Connectivity is really important, Chains are not important since local classifier does a good enough job.

An example - a graph from GraphViz showing what was predicted by local and how global fixed it - it should have both connectivity and triad issues.

5.4 Full pipeline

If we have this we can briefly explain about our first step system and show some results. This is good to say we do everything and bad if this really sucks.

6 Conclusion

In this paper we presented the task of process extraction and a method for extracting processes. We focused on extracting relations between event triggers. We also release publicly a data set for the scientific community. We have shown that by taking advantage of the global structure of a process we can improve performance.

Future work - adding more constraints - Mengqiu's idea. This may results in inference problems (it does) and so we can try think of smarter inference. There is the problem of very little data and we can think about using data from other domains and do adaptations. We want to do the full pipeline jointly.

References

- James F. Allen. 1983. Maintaining knowledge about temporal intervals. *Commun. ACM*, 26(11):832–843.
- Jari Björne, Juho Heimonen, Filip Ginter, Antti Airola, Tapio Pahikkala, and Tapio Salakoski. 2011. Extracting contextualized complex biological events with rich graph-based feature sets. *Computational Intelligence*, 27(4):541–557.
- Neil Campbell and Jane Reece. 2005. *Biology*. Benjamin Cummings.
- Nathanael Chambers and Daniel Jurafsky. 2008. Jointly combining implicit constraints improves temporal ordering. In *EMNLP*.
- Marie-Catherine de Marneffe, Bill MacCartney, and Christopher D. Manning. 2006. Accurate unlexicalized parsing. In *LREC*.
- Pascal Denis and Philippe Muller. 2011. Predicting globally-coherent temporal structures from texts via endpoint inference and graph decomposition. In *IJCAI*.
- Quang Do, Wei Lu, and Dan Roth. 2012. Joint inference for event timeline construction. In *EMNLP-CoNLL*, pages 677–687.
- Christiane Fellbaum, editor. 1998. MIT Press.
- Jenny Rose Finkel and Christopher D. Manning. 2008. Enforcing transitivity in coreference resolution. In *ACL (Short Papers)*, pages 45–48.
- Jin-Dong Kim, Tomoko Ohta, Sampo Pyysalo, Yoshinobu Kano, and Junichi Tsujii. 2009. Overview of bionlp 09 shared task on event extraction. In *Proceedings of BioNLP*.
- Jin-Dong Kim, Sampo Pyysalo, Tomoko Ohta, Robert Bossy, and Junichi Tsujii. 2011. Overview of bionlp shared task 2011. In *Proceedings of BioNLP*.
- Dan Klein and Christopher D. Manning. 2003. Accurate unlexicalized parsing. In *ACL*.
- André L. Martins, Noah A. Smith, and Eric P. Xing. 2009. Concise integer linear programming formulations for dependency parsing. In *ACL/IJCNLP*, pages 342–350.
- David McClosky and Christopher D. Manning. 2012. Learning constraints for consistent timeline extraction. In *EMNLP-CoNLL*, pages 873–882.
- David McClosky, Mihai Surdeanu, and Christopher D. Manning. 2011. Event extraction as dependency parsing. In *ACL*, pages 1626–1635.
- Makoto Miwa, Rune Sætre, Jin-Dong Kim, and Jun'ichi Tsujii. 2010. Event extraction with complex event classification using rich features. *J. Bioinformatics and Computational Biology*, 8(1).
- Hoifung Poon and Lucy Vanderwende. 2010. Joint inference for knowledge extraction from biomedical literature. In *HLT-NAACL*, pages 813–821.
- James Pustejovsky, José M. Castaño, Robert Ingria, Roser Sauri, Robert J. Gaizauskas, Andrea Setzer, Graham Katz, and Dragomir R. Radev. 2003. Timeml: Robust specification of event and temporal expressions in text. In *New Directions in Question Answering*, pages 28–34.
- Roi Reichart and Regina Barzilay. 2012. Multi-event extraction guided by global constraints. In *HLT-NAACL*, pages 70–79.

- Sebastian Riedel and Andrew McCallum. 2011. Fast and robust joint models for biomedical event extraction. In *Proceedings of the Conference on Empirical methods in natural language processing (EMNLP '11)*.
- Kristina Toutanova, Aria Haghighi, and Christopher D. Manning. 2008. A global joint model for semantic role labeling. *Computational Linguistics*, 34(2):161–191.
- Katsumasa Yoshikawa, Sebastian Riedel, Masayuki Asahara, and Yuji Matsumoto. 2009. Jointly identifying temporal relations with markov logic. In *ACL/IJCNLP*.